

CSE450 Capstone Project Report

ATLAS: An Adaptive Tutoring System using Bayesian Knowledge Tracing over a Prerequisite Skill Graph

Submitted by

Sakif Naieb Raiyan	2105065
Sayaad Muzahid Masfi	2105066
Aurchi Chowdhury	2105083
Omar Abdur Razzaque	2105094
Rubaiyat - E- Zaman	2105104

Supervised by

Dr. Md. Shariful Islam Bhuyan
Associate Professor, Department of CSE, BUET



Department of Computer Science and Engineering
Bangladesh University of Engineering and Technology
Dhaka, Bangladesh

Session 2026–2027
September 2026

CANDIDATES' DECLARATION

We declare that the work presented in this report, titled “ATLAS: An Adaptive Tutoring System using Bayesian Knowledge Tracing over a Prerequisite Skill Graph”, is the outcome of the engineering work carried out by us under the supervision named on the title page.

All prior work, third-party code, libraries, data, copyrighted material, and tools used in the project have been properly acknowledged. Any use of generative AI tools has been disclosed. We accept accountability and responsibility for the accuracy and integrity of this report and for our individual contributions to the project.

Sakif Naieb Raiyan (2105065)

Sayaad Muzahid Masfi (2105066)

Aurchi Chowdhury (2105083)

Omar Abdur Razzaque (2105094)

Rubaiyat - E- Zaman (2105104)

ACKNOWLEDGEMENT

We thank our supervisor for the guidance and critical questioning that shaped this project, and in particular for pressing us on the point that ultimately defined the work: that an adaptive tutor is only as good as the knowledge structure underneath it, and that a shallow skill map cannot be rescued by a sophisticated inference algorithm sitting on top of it.

We are grateful to the Department of Computer Science and Engineering, BUET, for the CSE 450 Capstone framework, which gave this work both its scope and its deadlines.

The subject-matter content of this system is drawn from standard Higher Secondary Certificate (HSC) textbooks and past admission-test question banks for Bangladeshi engineering universities. We acknowledge the authors and publishers of that material; it is used here strictly as source material for an academic prototype and is not redistributed as part of the system.

ABSTRACT

Every year a large cohort of Bangladeshi students prepares for the engineering university admission tests (BUET, KUET, RUET, CUET) by working through undifferentiated question banks. A student who repeatedly fails questions on integration by parts is rarely told that the actual gap is algebraic manipulation two levels below; coaching centres cannot diagnose at that granularity at scale, and existing digital platforms largely serve the same fixed content to everyone. The engineering problem is therefore one of inference under uncertainty: from a sparse, noisy sequence of multiple-choice responses, locate a learner’s knowledge frontier precisely enough to act on it.

This report presents ATLAS, an adaptive tutoring system that models a learner’s latent mastery with Bayesian Knowledge Tracing (BKT) over a directed acyclic graph (DAG) of prerequisite relations between granular skills. Three extensions to textbook BKT carry the design. First, the guess and slip parameters are conditioned on the Bloom’s Taxonomy level of the question, so that a correct answer on a recall item and a correct answer on an analysis item update belief by different amounts. Second, a correct response propagates upward through the prerequisite graph, raising every ancestor to at least the descendant’s posterior, since demonstrating a dependent skill is evidence for the skills it depends on. Third, a skill’s learning-rate parameter is gated conjunctively: it rises only when *all* of its direct prerequisites are mastered, preventing a learner from appearing to advance past an unmastered foundation. An adaptive diagnostic selects each next question by a balanced-volume heuristic over untested ancestors and descendants, at one Bloom level above the learner’s current estimated band.

The system was built as a FastAPI backend, a React frontend rendering Bangla text with embedded $\text{\LaTeX}$ mathematics, and a PostgreSQL (Supabase) store. A retrieval-augmented generation pipeline over six HSC textbooks (4,474 parsed question records) produces the question bank, with a two-model generate-then-verify design in which a second model blind-re-solves each generated item before it is accepted. The knowledge structure itself was the dominant engineering effort: an initial 430-skill ontology derived from a sample of the corpus proved too sparse to support the gating policy, and was rebuilt from the full corpus into 1,688 skills and 1,743 prerequisite edges covering all 183 syllabus topics, validated as an acyclic graph of depth 9 with no isolated skills.

Verification is structural and functional rather than empirical: an 18-check end-to-end harness exercises the diagnostic, mastery propagation, the ancestor-ordering invariant and prerequisite spillover against an in-memory database stand-in, and passes completely; the ontology passes an automated structural validator with zero errors. The system has not yet been evaluated with human learners, and establishing that its mastery estimates predict real performance remains the principal item of future work.

Keywords: Bayesian Knowledge Tracing; adaptive learning; prerequisite graph; Bloom’s Taxonomy; intelligent tutoring systems; retrieval-augmented question generation

Contents

Candidates' Declaration	i
Acknowledgement	ii
Abstract	iii
1 Problem Identification and Formulation	1
1.1 Background and Motivation	1
1.2 Problem Statement	2
1.3 Complexity of the Problem	3
1.4 Project Objectives and Scope	4
1.4.1 Objectives	4
1.4.2 Scope	5
2 Preliminary Research	6
2.1 Literature Review	6
2.1.1 Modelling what a learner knows	6
2.1.2 Cognitive level as a modelling variable	7
2.1.3 Automatic generation of assessment items	7
2.1.4 What the literature establishes, and what it leaves open	7
2.2 Existing Solutions and Technologies	8
2.3 Market and User Research	9
2.4 Feasibility Study	9
3 Requirements Analysis	11
3.1 Project Objectives and Constraints	11

3.1.1	Stakeholders and their expectations	11
3.1.2	Functional requirements	13
3.1.3	Non-functional requirements	13
3.1.4	Constraints	13
3.2	Regulatory Requirements and Standards	15
3.3	Formulation of Solutions	15
4	System Design and Implementation	17
4.1	Functional Design and Sub-problem Partitioning	17
4.2	Design Refinement and Simulation	19
4.2.1	The knowledge-tracing model	19
4.2.2	Conditioning the noise parameters on cognitive level	20
4.2.3	The three-phase mastery update policy	20
4.2.4	Adaptive question selection	21
4.2.5	Prerequisite spillover	22
4.2.6	The skill graph, and its rebuild	22
4.2.7	Data model	23
4.2.8	API surface	23
4.2.9	Content generation pipeline	23
4.2.10	Design calculations	25
4.3	Application of Tools	25
4.4	Implementation	26
4.5	Deployment	27
5	Investigation and Evaluation	29
5.1	Experimental Investigation	29
5.1.1	Investigation 1: is local model inference viable for content generation? .	29
5.1.2	Investigation 2: is the ontology's granularity sufficient to support the gating policy?	30

5.1.3	Investigation 3: can prerequisite edges be proposed mechanically? . . .	30
5.1.4	Investigation 4: do automated quality heuristics over-flag?	31
5.2	Performance Evaluation	32
5.2.1	Setup and acceptance criteria	32
5.2.2	Results	32
5.2.3	Evaluation against requirements	32
5.2.4	Evaluation against objectives	32
5.2.5	Limitations of this evaluation	32
5.3	Deviations and Design Revisions	35
5.3.1	Topic filter matched no rows (found, analysed, fixed)	35
5.3.2	Session state is not persistent	36
5.3.3	Authentication is username-only	36
5.3.4	The rebuilt ontology is not yet the served ontology	36
6	Project Management, Teamwork, and Impact	37
6.1	Project Management and Finance	37
6.1.1	Planning and Risk Management	37
6.1.2	Budgeting and Resource Identification	39
6.1.3	Economic Analysis and Financial Prospecting	39
6.1.4	Sustainability in Business and Commercialisation	41
6.2	Individual and Teamwork	41
6.2.1	Participation and Contribution	41
6.2.2	Collaboration and Conflict Resolution	42
6.2.3	Leadership and Direction	42
6.2.4	Multidisciplinary Engagement	43
6.3	Impact on Society	43
6.4	Environmental Impact and Life Cycle Sustainability	44
6.5	Ethics	45
6.5.1	Equity and Inclusivity	45

6.5.2	Accountability and Personal Responsibility	45
6.5.3	Proper Use of Intellectual Property	46
6.5.4	Professionalism and Ethical Codes	46
7	Conclusion	48
7.1	Conclusion and Future Work	48
7.1.1	Achievement against objectives	48
7.1.2	Limitations	49
7.1.3	Future work	49
7.2	Continuous Engagement and Self-directed Learning	50
	References	53

List of Figures

4.1	System architecture. Solid boxes are the online serving path; dashed boxes are offline processes that run before serving and write their artefacts into the compiled ontology and the question bank.	17
4.2	A fragment of the skill graph, illustrating the granularity target. “Integration by parts” is not one node but a chain: recognising applicability, choosing the decomposition, and executing it — each separately testable, and each with its own prerequisites, so that a failure can be localised to the step that actually broke.	22
6.1	Project schedule. The ontology rebuild in September was not in the original plan; it was forced by Investigation 2 (Section 5.1).	39

List of Tables

1.1	Complex engineering problem characteristics (P1–P7) and evidence from this project	3
1.2	Knowledge profiles referenced by P1	4
2.1	Existing solutions surveyed, and their limitations for this problem	8
3.1	Stakeholders, expectations, and conflicts	11
3.2	Functional requirements	12
3.3	Non-functional requirements	13
3.4	Standards and regulatory requirements, and how the design responds	14
3.5	Evaluation of candidate solutions against criteria drawn from the requirements .	15
4.1	Subsystem responsibilities	18
4.2	Bloom-conditioned guess and slip parameters, and the mastery band each level occupies	20
4.3	Legacy ontology versus full-corpus rebuild	23
4.4	Rebuilt ontology by subject	23
4.5	Database schema	24
4.6	API endpoints	24
4.7	Tool selection, justification, and relevant limitations	25
5.1	Structural effect of the ontology rebuild and its editorial pass	30
5.2	End-to-end verification harness: checks and what each establishes	33
5.3	Requirements traceability and status	34
6.1	Risk register: principal risks, their realisation, and mitigation	38
6.2	Development resources: actual and notional cost	40

6.3	Indicative operating cost for a production deployment (order-of-magnitude estimate, not a quotation)	40
6.4	Contribution by subsystem	42
7.1	Achievement against the objectives of Section 1.4	48

Chapter 1

Problem Identification and Formulation

1.1 Background and Motivation

Admission to the engineering universities of Bangladesh — BUET, KUET, RUET and CUET — is decided by a single competitive examination taken after the Higher Secondary Certificate (HSC). The test covers Mathematics, Physics and Chemistry, and the ratio of candidates to available seats is severe enough that preparation is effectively a year-long full-time activity for those who attempt it. The dominant preparation method is volume: a student buys one or more question banks compiled from previous years' papers and works through them, usually alongside coaching-centre classes that move at a fixed pace through a fixed syllabus.

This method has a structural weakness that has nothing to do with effort. When a student fails a question, the question bank tells them the correct answer; it does not tell them *why* they failed. A wrong answer on a problem requiring integration by parts may indicate that the student does not recognise when the technique applies, or cannot choose u and dv correctly, or chose them correctly but made an algebraic error in the resulting simpler integral, or does not know the derivative of the function involved. These are four different deficits requiring four different remedies, and they are invisible in a score. A human tutor localises them by watching the student work — which is precisely what one-to-one tutoring is effective at, and precisely what does not scale to a national cohort [1].

The consequence is misallocated study time. A student who does not know which of their foundations is weak has no rational basis for choosing what to study next, and defaults to either linear syllabus coverage or to practising the topics where failure is most visible — often the topics *downstream* of the actual gap. Time spent re-attempting integration problems does not repair an algebra deficit; it merely produces repeated failure on the same prerequisite, which is both inefficient and demoralising.

What is needed is a system that infers, from the student's own answers, where in the structure of the subject their knowledge stops — and then selects the next question at that frontier rather than at the point in the syllabus the calendar happens to have reached.

1.2 Problem Statement

Problem. Given a learner and a sequence of their responses to multiple-choice questions, estimate the learner’s mastery of every skill in a subject domain, and select the next question so as to localise and repair the learner’s knowledge frontier.

Formulated in engineering terms:

- **Entities.** A set of *skills* $S = \{s_1, \dots, s_n\}$, each an atomic, independently testable unit of knowledge or procedure, annotated with one Bloom’s Taxonomy cognitive level. A set of directed *prerequisite* edges $E \subseteq S \times S$, where $(s_i, s_j) \in E$ means s_i must be held before s_j is reachable. A bank of questions, each tagged with the skill it tests, its Bloom level and its topic. A set of learners.
- **Inputs.** For a learner u , an observed sequence of binary outcomes $(q_1, c_1), (q_2, c_2), \dots$ where q_t is the question served at step t and $c_t \in \{\text{correct}, \text{incorrect}\}$.
- **Latent state.** For each skill, $P(L_{u,s}) \in [0, 1]$, the probability that learner u has learned skill s . This is never observed directly; it must be inferred from noisy evidence.
- **Outputs.** (i) The posterior $P(L_{u,s})$ for all $s \in S$ after each response; (ii) the identity of the next question to serve.
- **Quantities to be guaranteed.** The graph (S, E) must be acyclic, or the notion of “prerequisite” is incoherent and propagation does not terminate. Mastery estimates must respect the graph: no skill may be credited as mastered while an ancestor it depends on is estimated as unmastered. Question selection must terminate even when the bank has no item matching the ideal specification.
- **Assumptions.** Responses are conditionally independent given the latent state; a learner does not forget within a session (no decay term); each question tests primarily one skill; the prerequisite structure is the same for all learners.

Engineering knowledge demanded. The formulation above is not solvable by routine application of standard components. It requires: probability theory and Bayesian inference (the latent-state update); graph theory (acyclicity checking, topological ordering, transitive reduction, ancestor traversal); algorithm design (the question-selection heuristic and its fallback ladder); database design under referential constraints; the pedagogical literature on knowledge tracing and cognitive taxonomies, which supplies the model family and the level annotation scheme; and, for the content pipeline, applied natural language processing — retrieval-augmented generation, embedding-based retrieval, and the validation of machine-generated assessment items.

1.3 Complexity of the Problem

A complex engineering problem must have characteristic P1 and some or all of P2–P7 (Table 1.1). P1 requires in-depth knowledge at the level of one or more of the knowledge profiles K3–K8 listed in Table 1.2.

Table 1.1: Complex engineering problem characteristics (P1–P7) and evidence from this project

Code	Definition	Evidence from this project
P1	Depth of knowledge required. Cannot be resolved without in-depth engineering knowledge at the level of one or more of K3, K4, K5, K6, or K8, which allows a fundamentals-based, first-principles analytical approach.	K3, K4, K5, K8. K3: the Bayesian update is derived from first principles, not called from a library — the posterior and transition equations are implemented directly (Section 4.2.1). K4: knowledge tracing and cognitive-taxonomy literature supplies the model family [2,3]. K5: the DAG, its acyclicity guarantee and transitive reduction are what make the gating policy implementable. K8: the design engages directly with the research literature on BKT variants [4–6] and on prerequisite-structured student models [7].
P2	Range of conflicting requirements. Involves wide-ranging or conflicting technical, engineering, and other issues.	Granularity conflicts with question-bank coverage: finer skills localise gaps better but multiply the items needed (the rebuilt ontology raised the generation space from 430×6 to $1,688 \times 6$ tuples). Diagnostic length conflicts with learner patience: 30 questions cannot cover 430 skills, forcing inference by propagation rather than measurement. Adaptivity conflicts with stability: an aggressive learning rate tracks change quickly but over-reacts to a lucky guess. Section 3.3 resolves these explicitly.
P3	Depth of analysis required. Has no obvious solution and requires abstract thinking and originality in analysis to formulate suitable models.	Textbook BKT is per-skill and independent; it cannot express “mastery of s_j is impossible while s_i is unmastered.” Making it graph-aware required an original three-phase policy (Bloom-conditioned update, ancestor pull-up, conjunctive transition gating) that is not a published algorithm but a composition designed for this problem (Section 4.2.3).
P4	Familiarity of issues. Involves infrequently encountered issues.	Bangla-language STEM content with embedded L ^A T _E X mathematics is outside the well-trodden path: it broke JSON escaping (a corruption affecting a measured 13.4% of one generated batch), broke text rendering in the browser, and is poorly served by English-centric embedding models.

continued on next page

Table 1.1 — continued from previous page

Code	Definition	Evidence from this project
P5	Extent of applicable codes. Is outside problems encompassed by standards and codes of practice for professional engineering.	No standard governs what a “skill” is, how finely a syllabus should be decomposed, or when an automatically generated assessment item is fit to serve to a student. These decisions were made against explicit project-defined criteria (Section 3.2) because no code of practice supplies them.
P6	Extent of stakeholder involvement and conflicting requirements. Involves diverse groups of stakeholders with widely varying needs.	Learners want the shortest path to a higher score; teachers want syllabus fidelity and correct content; the system’s own policy wants maximal information per question. These pull in different directions — the most informative question is often not the most encouraging one. Section 3.1 tabulates the conflicts.
P7	Interdependence. Is a high-level problem including many component parts or sub-problems.	Five interdependent subsystems: the knowledge ontology, the inference engine, the question bank and its generation pipeline, the serving API, and the learner-facing interface. The ontology is upstream of all others — its granularity determines what the engine can infer and what the generator must produce, which is why a sparse first ontology forced a full rebuild (Section 4.2.6).

Table 1.2: Knowledge profiles referenced by P1

Code	Description
K3	A systematic, theory-based formulation of engineering fundamentals required in the engineering discipline.
K4	Engineering specialist knowledge that provides theoretical frameworks and bodies of knowledge for the accepted practice areas in the engineering discipline; much is at the forefront of the discipline.
K5	Knowledge that supports engineering design in a practice area.
K6	Knowledge of engineering practice (technology) in the practice areas in the engineering discipline.
K8	Engagement with selected knowledge in the research literature of the discipline.

1.4 Project Objectives and Scope

1.4.1 Objectives

The project set the following verifiable objectives. Chapter 5 evaluates the delivered system against each.

- O1. Construct a prerequisite skill graph** for the HSC Mathematics, Physics and Chemistry syllabus at a granularity fine enough that a single node names one independently testable action, and verify that the graph is acyclic and free of isolated nodes.
- O2. Implement a knowledge-tracing engine** that maintains a per-learner, per-skill mastery estimate, updates it from each response by Bayesian inference, and conditions that inference on the cognitive level of the question.
- O3. Make the estimate graph-aware**, such that evidence propagates to prerequisite skills and no skill can be credited above an unmastered prerequisite.
- O4. Deliver an adaptive diagnostic** that localises a learner's frontier within a bounded number of questions (target 30) rather than testing every skill.
- O5. Deliver targeted practice** that drives a chosen topic to a defined mastery threshold and detects when repeated failure is caused by a missing prerequisite rather than by the topic itself.
- O6. Produce a question bank** aligned to the skill graph, in Bangla with correct mathematical typesetting, via an automated pipeline with a verification stage.
- O7. Deliver a usable web application** exposing the diagnostic, practice and a mastery view to a learner.
- O8. Verify the engine end to end** with an automated test harness that runs without external credentials.

1.4.2 Scope

Included: the three admission-test subjects; multiple-choice items with four options; a single-process web deployment sufficient for demonstration; Bangla content with $\text{\LaTeX}$ mathematics; English and Bangla interface text.

Explicitly excluded: written/subjective answer grading; handwriting or image input; mobile-native applications; teacher-facing classroom analytics; multi-tenant institutional deployment; live proctoring; and any claim of improved examination outcomes, which would require a controlled study with human participants that this project did not undertake (Section 5.2.5).

Chapter 2

Preliminary Research

2.1 Literature Review

2.1.1 Modelling what a learner knows

The problem of inferring an unobservable knowledge state from observable responses has a long history in intelligent tutoring systems. Bayesian Knowledge Tracing (BKT), introduced by Corbett and Anderson [2], remains the reference formulation: a skill’s mastery is a hidden binary variable, and four parameters govern its evolution — a prior $P(L_0)$, a learning rate $P(T)$, and two noise terms, the slip probability $P(S)$ (knowing but answering wrongly) and the guess probability $P(G)$ (not knowing but answering correctly). The model’s durability comes from the fact that these four parameters are interpretable to a teacher, which is not true of most of its successors.

The literature has extended BKT along three axes that are directly relevant here. The first is *parameter individualisation*: Yudelson et al. [5] fit per-student parameters and report improved predictive accuracy, establishing that treating all learners as identical is a real source of error. The second is *contextualisation of noise*: Baker et al. [4] estimate slip and guess per response rather than per skill, on the argument that a guess on a four-option item is not the same event as a guess on an open response. Pardos and Heffernan [6] make the related argument for item difficulty. The third is *structure between skills*: Käser et al. [7] model students with dynamic Bayesian networks in which skills are linked by prerequisite relations, and report that exploiting that structure improves prediction over independent per-skill models.

Competing model families exist. Learning Factors Analysis [8] and Performance Factors Analysis [9] replace the hidden-state formulation with logistic regression over counts of prior successes and failures. Deep Knowledge Tracing [10] replaces it with a recurrent neural network and reports substantial accuracy gains on large datasets. These are relevant as alternatives considered (Section 3.3) but share a property that disqualified them here: their internal state is not interpretable as “this student has not mastered this specific skill,” which is exactly the output this system must produce in order to act on it.

A separate and older tradition, knowledge space theory [11, 12], formalises a domain as a partially ordered set of knowledge states and asks which questions most efficiently locate a learner within it. This is the theoretical ancestor of the prerequisite graph used here, and it supplies the central insight that a well-structured domain permits assessment far shorter than the number of skills, because observing one response constrains many states at once.

2.1.2 Cognitive level as a modelling variable

Bloom’s taxonomy [13], in its revised form [3], classifies learning objectives into six ascending cognitive levels from Remember to Create. Its use in tutoring systems is usually editorial — a way of labelling items — rather than a variable inside the student model. The observation this project builds on is that the taxonomy has a direct probabilistic reading: a four-option recall question can be answered correctly by chance far more easily than an analysis question, and a learner who genuinely knows an analysis-level skill is more likely to slip on it through a procedural error. That makes the Bloom level a natural conditioning variable for $P(G)$ and $P(S)$, which is the refinement Baker et al. [4] argue for on empirical grounds but implemented from a pedagogical prior rather than a fitted one.

Vygotsky’s zone of proximal development [14] supplies the selection rule: instruction is most effective just beyond what the learner can already do unaided. Translated into this system’s terms, the next question should sit one Bloom level above the learner’s current estimated band — which is the rule the diagnostic implements.

2.1.3 Automatic generation of assessment items

Kurdi et al. [15] survey automatic question generation for education and report the field’s persistent difficulty: generated items are often syntactically fine but pedagogically hollow, and evaluation of their quality is usually thin. Retrieval-augmented generation [16] addresses part of this by grounding generation in retrieved source material rather than in model memory, which both improves factual accuracy and ties items to a specific curriculum. Multilingual sentence embeddings [17] make such retrieval feasible over a non-English corpus.

2.1.4 What the literature establishes, and what it leaves open

Established: BKT is an adequate and interpretable model of per-skill mastery; its noise parameters should not be uniform across items; prerequisite structure between skills carries real predictive value; and short adaptive assessment is possible when the domain is well structured.

Left open, and therefore the space this project works in: the literature provides no standard

method for *constructing* a fine-grained prerequisite graph for a specific national syllabus — the graphs in published work are either small, hand-built by domain experts, or learned from response data that does not exist before a system is deployed. Nor does it settle how a Bloom-conditioned, graph-propagating variant of BKT should behave when the question bank is incomplete, which is the normal condition of a system built before it has users.

2.2 Existing Solutions and Technologies

Table 2.1: Existing solutions surveyed, and their limitations for this problem

Category	Examples	Strengths and limitations for this problem
Adaptive commercial platforms	ALEKS; Knewton-style adaptive engines	Strength: mature adaptive engines with knowledge-space foundations and demonstrated scale. Limitation: closed, licensed per seat, and built around Western curricula — no HSC syllabus coverage, no Bangla content, and no ability to add either.
Bangladeshi ed-tech platforms	10 Minute School, Shikho and similar	Strength: strong local content, Bangla-first, priced for the local market, already trusted by the target cohort. Limitation: predominantly content-delivery and fixed mock tests; assessment reports a score per test or per chapter rather than inferring a per-skill knowledge state, so the student still has to decide what to study next.
Generic quiz and spaced-repetition tools	Quizlet, Anki, Google Forms	Strength: free, simple, and effective for pure recall [18]. Limitation: no model of prerequisite structure at all — a card is either due or not due — so they cannot diagnose <i>why</i> an item was failed or redirect to an upstream deficit.
Printed question banks and coaching centres	The incumbent method	Strength: authoritative content, human explanation, peer environment. Limitation: fixed pace for a heterogeneous room; diagnosis at the granularity of a chapter, not a skill; and the marginal cost of one-to-one attention is precisely what prevents it scaling [1].
Open-source ITS frameworks	pyBKT and similar research toolkits	Strength: reference implementations of the model family, useful for validation. Limitation: libraries, not products — they supply the inference but neither the ontology, the content, nor the serving layer, which is where the majority of this project’s effort actually went.

The gap. No surveyed system combines (i) a fine-grained prerequisite model of *this* syllabus,

(ii) inference that is interpretable per skill and propagates across prerequisite structure, and (iii) Bangla content with correct mathematical typesetting. Platforms that have the content lack the model; platforms that have the model lack the content and the language. ATLAS targets that intersection.

2.3 Market and User Research

User base. The primary users are HSC-completing students in Bangladesh preparing for engineering admission tests, typically in a concentrated period between the HSC examination and the admission tests. Secondary users are repeat candidates and students in the HSC years using the system for syllabus practice rather than test preparation.

What users do instead. Current practice, established from the team’s own experience of the same examination and from the structure of the material that dominates the market, is: solve past-paper compilations; attend coaching classes; sit periodic mock tests that return a rank and a score. The gap this leaves is diagnostic, not informational — students do not lack content, they lack a reliable answer to “what should I study next?”

Demand characteristics. Willingness to spend on preparation is demonstrably high, since coaching enrolment is close to universal among serious candidates, but willingness to spend on *software* is low; the realistic price point for a digital tool is at or near zero for the student, with revenue arriving through institutions if at all (Section 6.1.3).

Evidence and its limits. The characterisation above is grounded in the team’s direct familiarity with the examination and an informal survey of available platforms and printed material. It is *not* grounded in a structured survey of prospective users with a defined sample; no such study was conducted. This is a genuine weakness of the project’s user research, stated here rather than obscured, and Section 7.1 lists it as future work.

2.4 Feasibility Study

Technical. Feasible, and demonstrated. The inference model is computationally trivial (a constant-time update per response plus an ancestor traversal bounded by graph depth, measured at 9). The content pipeline was the real technical risk: it depends on a large language model producing correct Bangla mathematical content at scale. That risk was retired empirically — an initial attempt using local models on CPU hardware proved infeasible on throughput grounds and was abandoned, and a hosted API with a generate-then-verify design was substituted successfully (Section 4.2.9).

Economic. Feasible at prototype scale at essentially zero marginal cost: every service used (database, model API, hosting) has a free tier adequate for development and demonstration. It is *not* feasible to assume those free tiers extend to production; Section 6.1.3 models the real operating cost.

Operational. The adoption question is the binding one. The system asks a student to spend 30 questions on a diagnostic before it gives them anything, which is a real cost to a user under time pressure who is accustomed to immediate practice. Mitigation is to make the diagnostic itself feel like practice — it serves genuine past-paper-style questions, not artificial calibration items.

Legal and regulatory. Feasible with attention. The system stores learner performance data, which is personal data; the source content is drawn from copyrighted textbooks and question banks, which constrains redistribution. Both are addressed in Section 3.2 and Section 6.5.3.

Conclusion. The project is feasible, and the principal assumption it depends on is that a machine-constructed skill ontology, validated structurally but not yet by domain experts, is an adequate foundation for the inference built on top of it. That assumption is explicitly flagged as unretired (Section 5.2.5).

Chapter 3

Requirements Analysis

3.1 Project Objectives and Constraints

3.1.1 Stakeholders and their expectations

Table 3.1: Stakeholders, expectations, and conflicts

Stakeholder	Expectation	Conflict with other stakeholders
Admission candidate (primary user)	Shortest path to a higher score; immediate practice; questions that feel like the real examination; no wasted time.	Wants to start practising at once; the system wants a 30-question diagnostic first. Wants encouragement; the most informative question is the one they are about 50–70% likely to fail.
Teacher / coaching instructor	Syllabus fidelity, factually correct content, results that map to chapters they teach.	Thinks in chapters; the engine thinks in skills that cut across chapters. Requires content correctness guarantees that a generative pipeline cannot fully give.
Project team (developer)	A system that is demonstrably correct, maintainable, and finishable within one academic session.	Correctness argues for hand-authored content and expert-reviewed ontology; the schedule argues for automation.
Supervisor / examiner	Evidence of engineering depth, rigour, and honest evaluation.	An honest evaluation reports unvalidated components, which conflicts with the team’s interest in presenting a finished product.
Content rights holders	Their material not redistributed.	The system needs that material as source for retrieval and generation.

Table 3.2: Functional requirements

ID	Requirement	Objective
FR1	A learner can create an account and sign in.	O7
FR2	The system presents a catalogue of courses, sections and topics derived from the skill ontology.	O1, O7
FR3	A learner can start a diagnostic for a section; the system serves questions one at a time, adaptively selected.	O4
FR4	Each response updates the learner's mastery estimate for the tested skill using Bayesian inference conditioned on the question's Bloom level.	O2
FR5	A correct response propagates upward, raising ancestor skills to at least the descendant's posterior.	O3
FR6	A skill's learning rate is gated on all of its direct prerequisites being mastered.	O3
FR7	Mastery views are locked until the section diagnostic is complete.	O4
FR8	After the diagnostic, a learner can practise a chosen topic until every skill in it reaches the mastery threshold.	O5
FR9	Repeated failure attributable to a missing prerequisite triggers injection of prerequisite-focused questions.	O5
FR10	A learner can view their mastery as both a graph and a table, with subject and topic labels.	O7
FR11	A learner can reset a section and retake its diagnostic.	O7
FR12	Questions render Bangla text with embedded $\text{\LaTeX}$ mathematics correctly.	O6, O7
FR13	The question bank is generated from the source corpus and aligned to skill, Bloom level and topic.	O6
FR14	The interface is available in English and Bangla.	O7

Table 3.3: Non-functional requirements

ID	Attribute	Requirement
NFR1	Correctness	The prerequisite graph must be acyclic; mastery estimates must never place an ancestor below its descendant; these must be checked automatically, not assumed.
NFR2	Responsiveness	Next-question selection must complete fast enough to feel immediate in a browser session; the search must terminate even when no ideal item exists.
NFR3	Robustness	Absence of a matching question must degrade gracefully through a defined fallback ladder rather than failing the session.
NFR4	Maintainability	The ontology must be compiled from a single source of truth by a reproducible build, not hand-edited in its served form.
NFR5	Verifiability	The engine must be testable end to end without database credentials.
NFR6	Portability	The system must run on a standard developer machine with no GPU.
NFR7	Localisation	All learner-facing content must support Bangla, including mathematics.
NFR8	Security	Secrets must never enter version control; the database service key must remain server-side.

3.1.2 Functional requirements

3.1.3 Non-functional requirements

3.1.4 Constraints

Time: one academic session, with the team also carrying a full course load. **Budget:** effectively zero; every external service had to have a usable free tier (Section 6.1.2). **Hardware:** development on personal laptops with no GPU — this constraint alone eliminated the original local-model content-generation approach. **Data:** no existing learner response data was available, so no model parameter could be fitted empirically; all parameters had to be set from pedagogical priors. **Platform:** web, to avoid the cost of maintaining native clients. **Deployment:** a single process sufficient for demonstration, which is why session state is held in memory (Section 5.3.2).

Table 3.4: Standards and regulatory requirements, and how the design responds

Instrument	Why it applies	How the design complies
Data-protection obligations for personal data of minors and young adults [19]	The system stores per-learner performance data; many users are 17–19 years old.	Data minimisation: the schema stores only a username and per-skill mastery — no email, phone, address, date of birth or payment data. No third-party analytics or advertising trackers are embedded. Deletion is supported at section granularity by the retake endpoint. <i>Gap</i> : the current prototype has no password authentication, so this minimal data is not access-controlled; see Section 5.3.3.
OWASP Top 10 [20]	The system is a public-facing web application with a database behind it.	All database access goes through a parameterised client library rather than string-concatenated SQL, addressing injection. The privileged service key is held server-side in an environment file that is git-ignored and never shipped to the browser, addressing sensitive-data exposure. <i>Gap</i> : broken access control — username-only login and wide-open CORS are known violations retained deliberately for demonstration scope, and are the first items in the production hardening list (Section 7.1.3).
WCAG 2.1 [21]	The interface is the sole delivery channel and excludes users if inaccessible.	Semantic HTML controls, keyboard-operable question options, text-based mathematics rendered as accessible markup rather than images, and a bilingual interface reducing language exclusion. <i>Gap</i> : no formal audit against AA criteria was performed (Section 6.5.1).
Copyright in source textbooks and question banks	Content is derived from copyrighted printed material.	Source material is used as a private retrieval corpus, is not redistributed, and is excluded from the public repository’s distributable artefacts; generated items are new text grounded in that corpus rather than reproductions of it.
IEEE/ACM codes of ethics [22, 23]	The system makes inferences that steer a student’s study time during a high-stakes year.	Honest reporting of the system’s unvalidated components, and refusal to claim score improvement without evidence, follow directly from these codes (Section 6.5.4).

3.2 Regulatory Requirements and Standards

3.3 Formulation of Solutions

Three genuinely different architectures could each satisfy most of the functional requirements. They differ in where the intelligence lives.

Alternative A — Item Response Theory over a flat item pool. Model each *question* with a difficulty parameter and each learner with a single latent ability parameter; select the next item to maximise information about that ability. This is the classical computerised-adaptive-testing design.

Alternative B — Deep Knowledge Tracing. Model the learner’s entire response history with a recurrent network that predicts the probability of answering the next item correctly [10]; select items to maximise learning gain under the network’s predictions.

Alternative C — Bayesian Knowledge Tracing over a prerequisite DAG (selected). Model each skill separately with an interpretable four-parameter hidden-state model, link skills by prerequisite edges, and propagate evidence along those edges.

Table 3.5: Evaluation of candidate solutions against criteria drawn from the requirements

Criterion (source)	A: IRT	B: DKT	C: BKT+DAG
Produces an actionable per-skill diagnosis (FR4, and the core problem statement)	No — one ability scalar cannot say <i>which</i> skill is missing	No — hidden vector is not interpretable per skill	Yes
Can express prerequisite structure (FR5, FR6)	No	Only implicitly, if data suffices	Yes, explicitly
Trainable with zero historical learner data (project constraint)	Needs calibration data per item	Needs large response datasets	Yes — parameters set from pedagogical priors
Interpretable to a teacher (teacher stakeholder)	Partly	No	Yes
Computational cost on commodity hardware (NFR6)	Low	High (training)	Low
Robust when the question bank is incomplete (project reality)	Poor — needs calibrated items	Poor — needs coverage	Fair — fallback ladder plus propagation
Implementation effort within one session	Medium	High	Medium

Conclusion. Alternative C was selected. The decisive criteria are the first and third rows.

The project's entire premise is telling a student *which* foundation is missing, which A and B structurally cannot do: A compresses knowledge to one number, and B's state is a vector with no per-skill reading. Independently, both A and B require historical response data to calibrate, and the project began with none and would have none until it had users — a circular dependency that C avoids because its four parameters have defensible pedagogical defaults. B remains the most credible successor once real response data accumulates, and the engine was deliberately built so that the model can be replaced without touching its call sites (Section 4.2.1).

Chapter 4

System Design and Implementation

4.1 Functional Design and Sub-problem Partitioning

The problem decomposes into five subsystems with a deliberate dependency order. The knowledge ontology is upstream of everything: it defines the skills the engine reasons about, the topics the interface displays, and the specifications the content pipeline must generate against. The inference engine depends on the ontology but not on the content; the content pipeline depends on the ontology but not on the engine; the API composes engine and content; the interface depends only on the API. This ordering is what allowed the ontology to be rebuilt late in the project (Section 4.2.6) without rewriting the engine.




Figure 4.1: System architecture. Solid boxes are the online serving path; dashed boxes are offline processes that run before serving and write their artefacts into the compiled ontology and the question bank.

Where standards shaped the partition. Two boundaries in Figure 4.1 exist for reasons from Section 3.2 rather than for reasons of convenience. First, the privileged database key lives only

Table 4.1: Subsystem responsibilities

Subsystem	Location	Responsibility
Knowledge ontology	Ontology/, Backend/tree_ data/	Defines skills, their Bloom levels, their topics, and the prerequisite edges between them. Compiled from source into the served JSON artefacts by a build script that enforces acyclicity and applies transitive reduction.
Inference engine	Backend/bkt.py, bloom_taxonomy. py, mastery_ updater.py, skill.py, skill_tree.py	Maintains and updates per-skill mastery; owns the graph, its cycle check and its topological ordering.
Session logic	Backend/ diagnostic.py, session state in server.py	Runs a diagnostic or practice session: which skill to test next, which Bloom level, which question, when to stop.
Serving API	Backend/server. py	Exposes catalogue, session and mastery endpoints; composes engine and database.
Content pipeline	data-gen/	Parses the corpus, builds a retrieval index, generates and verifies questions, loads them to the database.
Client	Frontend/	Renders the catalogue, question flow and mastery views; handles bilingual display and mathematical typesetting.

in the API process: the client never holds a credential that can write to the database, which is why there is no direct client-to-database arrow even though the database vendor supports one. Second, the corpus sits strictly inside the offline pipeline and is never served: generated items reach the database, but the copyrighted source text never leaves the build environment.

Safety considerations. The system carries no physical-safety risk, but it does carry a decision-quality risk: a learner may reallocate study time during a high-stakes year on the strength of a wrong inference. The design response is conservatism in the gating policy — the threshold for declaring a prerequisite mastered is deliberately high (0.95), so the system prefers to keep a learner practising a foundation slightly too long over releasing them from it too early.

4.2 Design Refinement and Simulation

4.2.1 The knowledge-tracing model

Each skill's mastery is a latent binary variable; the engine maintains $P(L_n)$, the probability that the learner has learned the skill after n opportunities. Four parameters govern its evolution: the prior $P(L_0)$, the learning rate $P(T)$, the slip probability $P(S)$ and the guess probability $P(G)$ [2]. Each response triggers a two-step update.

Step 1, the Bayesian posterior given the observed response:

$$P(L_n \mid \text{correct}) = \frac{P(L_n) (1 - P(S))}{P(L_n) (1 - P(S)) + (1 - P(L_n)) P(G)} \quad (4.1)$$

$$P(L_n \mid \text{incorrect}) = \frac{P(L_n) P(S)}{P(L_n) P(S) + (1 - P(L_n)) (1 - P(G))} \quad (4.2)$$

Step 2, the learning transition — the chance that the opportunity itself taught the skill:

$$P(L_{n+1}) = P(L_n \mid \text{evidence}) + (1 - P(L_n \mid \text{evidence})) P(T) \quad (4.3)$$

Mastery is persisted as a percentage on $[0, 100]$ and related to the probability form by $P(L) = \text{mastery}/100$. The implementation is deliberately *stateless* — it takes a probability and an observation and returns a probability, holding no learner state — so that the entire model can be replaced by an individualised or neural variant [5, 10] without altering a single call site. That is the concrete design provision for the successor identified in Section 3.3.

4.2.2 Conditioning the noise parameters on cognitive level

The refinement that distinguishes this engine from textbook BKT is that $P(G)$ and $P(S)$ are not constants but functions of the question’s Bloom level. A four-option recall item is guessable; an analysis item is not. Conversely, a learner who knows an analysis-level procedure has more opportunities to slip while executing it. Table 4.2 gives the values used.

Table 4.2: Bloom-conditioned guess and slip parameters, and the mastery band each level occupies

Bloom level	$P(G)$	$P(S)$	Mastery band
Remember	0.30	0.05	0.00 – 16.67
Understand	0.25	0.08	16.67 – 33.33
Apply	0.18	0.12	33.33 – 50.00
Analyze	0.12	0.18	50.00 – 66.67
Evaluate	0.08	0.22	66.67 – 83.33
Create	0.05	0.25	83.33 – 100.00

The 0–100 mastery scale is partitioned into six equal bands of $100/6 \approx 16.67$ points, one per level, so that a mastery value maps directly to a cognitive level. This is what makes the selection rule expressible: the next question is requested one band above the learner’s current band, which is the zone-of-proximal-development rule [14] rendered arithmetically.

4.2.3 The three-phase mastery update policy

Every response runs three phases in order.

Phase 1 — Bloom-conditioned BKT. Look up $P(G)$, $P(S)$ for the question’s Bloom level, compute the effective transition rate (Phase 3), and apply Equations (4.1) to (4.3) to the answered skill.

Phase 2 — Ancestor pull-up. If the response was correct, traverse the skill’s ancestors breadth-first and raise each to at least the answered skill’s new posterior:

$$P(a) \leftarrow \max(P(a), P(s)) \quad \forall a \in \text{ancestors}(s) \quad (4.4)$$

The justification is evidential: correctly executing a skill is evidence that its prerequisites are held, because the skill is not executable without them. The asymmetry is deliberate — an *incorrect* response propagates nothing downward to prerequisites, because failure is ambiguous. It may indicate a missing prerequisite, but it may equally indicate a slip at the top level, and penalising every ancestor for one wrong answer would make the estimate unstable. Locating the responsible prerequisite is instead the job of the spillover rules, which use option-level evidence rather than a blanket penalty.

Phase 3 — Conjunctive transition gating. A skill’s learning rate is not constant. For a skill s with direct prerequisites $\text{parents}(s)$:

$$P(T)_s = \begin{cases} \tau_{\text{base}} & \text{if } \min_{p \in \text{parents}(s)} P(L_p) \geq \theta \\ \tau_{\text{locked}} & \text{otherwise} \end{cases} \quad (4.5)$$

with $\theta = 0.95$, $\tau_{\text{base}} = 0.10$ in normal practice, 0.20 during a diagnostic (where the aim is to move the estimate quickly), and $\tau_{\text{locked}} = 0.01$. The gate is *conjunctive* — every prerequisite must clear the threshold, not merely the average or the best one. This is the formal expression of the project’s central claim: a learner cannot be credited with advancing on a skill while a foundation beneath it is missing. Its correctness depends entirely on $\text{parents}(s)$ being the *direct* prerequisites, which is why the ontology build applies transitive reduction (Section 4.2.6); a redundant shortcut edge would let a skill unlock while its true intermediate prerequisite was still unmastered.

4.2.4 Adaptive question selection

A diagnostic cannot test 430 skills in 30 questions, so it must choose skills whose outcome is maximally informative about the rest of the graph. The engine scores each untested skill by a balanced-volume heuristic:

$$\text{score}(s) = \frac{u_{\uparrow}(s) + u_{\downarrow}(s) + 1}{1 + |u_{\uparrow}(s) - u_{\downarrow}(s)|} \quad (4.6)$$

where $u_{\uparrow}$ and $u_{\downarrow}$ are the counts of untested ancestors and untested descendants. The numerator rewards skills that touch many unknowns; the denominator rewards *balance*. A skill with 20 unknown descendants and no unknown ancestors is a root whose failure tells us little about what lies above it; a skill sitting between two large unknown regions bisects the graph, and its outcome constrains both directions. This is the knowledge-space insight [11] applied as a greedy heuristic.

Priors are initialised by depth: root skills start at 0.70 mastery and the deepest skills at 0.20, interpolated linearly between, encoding the assumption that foundational material is more likely already held.

Having chosen a skill and a target Bloom level, the engine searches for an item through a fallback ladder, because the ideal item frequently does not exist in a partially generated bank: (i) same skill, nearby Bloom level, matching topic; (ii) same skill, nearby Bloom level, any topic; (iii) same topic, graph-nearby skills ordered by distance, preferring the lowest-mastery candidate. “Nearby Bloom” means the offsets $[0, +1, -1, +2, -2, +3, -3]$ tried in that order. If no tier yields an item, the skill is marked tested and the session advances — the design decision that a missing question must never stall a session (NFR3).

4.2.5 Prerequisite spillover

During topic practice, two rules detect that failure is caused by something upstream rather than by the topic itself. Rule A fires when at least three of the last five wrong answers implicate the same missing prerequisite, evidence that comes from the `option_missing_prerequisites` table — each distractor is annotated with the skill whose absence would lead a learner to choose it. Rule B fires when rolling accuracy over the last eight attempts falls below 40% *and* some candidate prerequisite has mastery below 0.60; the weakest such prerequisite is chosen. Either rule injects two prerequisite-focused questions before returning to the topic traversal. Rule A is precise but needs distractor annotations; Rule B is a coarser safety net that works without them.

4.2.6 The skill graph, and its rebuild




Figure 4.2: A fragment of the skill graph, illustrating the granularity target. “Integration by parts” is not one node but a chain: recognising applicability, choosing the decomposition, and executing it — each separately testable, and each with its own prerequisites, so that a failure can be localised to the step that actually broke.

The first ontology was derived from a sample of the corpus and yielded 430 skills across 66 topics with 471 edges. In use it proved too coarse: a topic held on average fewer than four skills, so a failure inside a topic could not be localised further, and the conjunctive gate had too few edges to act on. The graph was therefore rebuilt from the complete corpus — all 4,474 parsed records of all six books — at a granularity where a node names one action. Table 4.3 compares the two.

The build is reproducible rather than hand-maintained: a compiler reads the ontology source, applies an editorial layer (topic aliases and display labels), checks the graph for cycles, applies transitive reduction, and emits the JSON artefacts the server loads. Transitive reduction is not cosmetic — as Equation (4.5) shows, a shortcut edge would let a skill unlock while its true intermediate prerequisite was unmastered.

Table 4.3: Legacy ontology versus full-corpus rebuild

Property	Legacy (served)	Rebuild
Skills	430	1,688
Prerequisite edges	471	1,743
Topics covered	66	183 of 183
Corpus coverage	sample	all 4,474 records
Maximum graph depth	7	9
Isolated skills	—	0
Unresolved prerequisite references	—	0
Connected components (largest)	—	87 (482)

Table 4.4: Rebuilt ontology by subject

Subject	Skills	Edges	Edges per skill	Topics
Chemistry	886	868	0.97	59
Mathematics	446	499	1.10	52
Physics	356	376	1.06	72
Total	1,688	1,743	1.03	183

4.2.7 Data model

Two deliberate omissions are worth stating. There is no `subject` column anywhere: `subject` is a property of the ontology, derived at read time and plumbed through the API, so that the database never holds a second copy of a fact the ontology owns. And the transition probability $P(T)_s$ is not persisted: it is recomputed from current parent mastery on every request, so it can never fall out of sync with the values it is derived from.

4.2.8 API surface

The API is a flat, session-oriented REST surface of 14 endpoints (Table 4.6). The two session families are deliberately symmetric — `start`, `status`, `next`, `answer` — so the client can drive either flow with one component.

4.2.9 Content generation pipeline

The pipeline turns the corpus into items aligned to the ontology. Source books are parsed into structured records; a multilingual embedding model [17] indexes them; for each (skill, Bloom level) specification the top- k relevant records are retrieved and supplied to a generator model as grounding [16].

Table 4.5: Database schema

Table	Contents and constraints
<code>users</code>	Learner identity: surrogate UUID, unique username.
<code>user_skill</code>	Per-learner, per-skill mastery on [0, 100]; composite primary key; cascades on learner or skill deletion.
<code>skills</code>	Skill identifier and description; description is NOT NULL.
<code>questions</code>	Stem, Bloom level, and topic. <code>topic</code> is a foreign key to <code>ontology_topics</code> , so it stores a topic <i>code</i> , never a display label.
<code>question_options</code>	Four options per question, label constrained to A–D, correctness flag, and a per-option explanation.
<code>option_missing_</code> <code>prerequisites</code>	Maps a distractor to the skill whose absence explains choosing it — the evidence source for spillover Rule A.
<code>ontology_topics</code> , <code>ontology_skill_topics</code> , <code>ontology_skill_edges</code>	The ontology mirrored into the database for referential integrity of questions and skills.

Table 4.6: API endpoints

Endpoint	Purpose
GET <code>/catalog</code>	Courses, sections and topics from the compiled ontology
GET <code>/users/{name}/</code> , POST <code>/users/</code>	Look up and create a learner
GET <code>.../sections/{id}/state</code>	Diagnostic gate and lock state for a section
GET <code>.../sections/{id}/mastery</code>	Mastery graph and table for visualisation
POST <code>.../diagnostic/retake</code>	Reset section mastery and clear its sessions
POST <code>/diagnostic/start</code>	Begin a diagnostic session
GET <code>/diagnostic/{id}/status</code> , <code>/next</code>	Session progress; fetch next selected question
POST <code>/diagnostic/{id}/answer</code>	Submit a response; triggers the three-phase update
POST <code>/topic-practice/start</code>	Begin practice on one topic
GET <code>/topic-practice/{id}/status</code> , <code>/next</code>	Progress; next question, including injected spillover items
POST <code>/topic-practice/{id}/answer</code>	Submit a response; evaluates spillover rules

The design decision that matters most is the second stage. A generated item is not accepted on the generator’s own word: a *different, cheaper* model independently checks that the item is on-topic for the claimed skill, and *blind re-solves* it — it is shown the question and options but not the marked answer — and the item is accepted only if its independent answer agrees with the marked one. Letting a model grade its own output biases towards false confidence; an independent judge with no sight of the expected answer is a real check. The cost asymmetry is deliberate: generation uses the stronger model because it is the harder task, verification uses the cheaper one because agreement checking is easier than authoring, and verification calls are batched across several specifications to reduce call count.

4.2.10 Design calculations

Generation space. With one item per (skill, Bloom level) pair, the legacy ontology implies $430 \times 6 = 2,580$ specifications; the rebuilt ontology implies $1,688 \times 6 = 10,128$. At the observed generation pattern this is the dominant cost driver in the project and the reason adoption of the rebuilt ontology is gated on a budgeted generation run (Section 6.1.2).

Storage. A question with four options and explanations occupies on the order of a few kilobytes. A complete bank at rebuild scale is therefore of the order of tens of megabytes — immaterial against any database tier. Learner data is smaller still: one row per learner per skill, so a cohort of 10,000 learners against 1,688 skills is under 2×10^7 rows of two identifiers and a float, which is well within commodity capacity.

Inference cost. One response costs one constant-time Bayesian update plus an ancestor traversal bounded by graph depth (9) and a successor refresh over direct children. The engine is not the scaling bottleneck; database round trips during question search are.

4.3 Application of Tools

Table 4.7: Tool selection, justification, and relevant limitations

Tool	Why selected	Relevant limitations
Python 3	The engine is mathematical and graph-oriented; Python keeps Equation (4.1)–Equation (4.5) readable enough to audit against the policy document.	Slower than a compiled language — immaterial here, since the update is constant-time.

continued on next page

Table 4.7 — continued from previous page

Tool	Why selected	Relevant limitations
FastAPI + Uvicorn	Typed request/response models via Pydantic catch malformed payloads at the boundary; automatic OpenAPI documentation made the API self-describing for the front-end developers.	The asynchronous model gives no benefit as used, since the engine is synchronous and CPU-light.
Supabase (PostgreSQL)	A managed relational store with a usable free tier. Relational integrity is not optional here — foreign keys are what stop a question referencing a skill that no longer exists.	The free tier pauses idle projects; the client library adds a network round trip where raw SQL would allow one query.
React 18 + Vite	Component model suits a question-at-a-time flow with shared state; Vite’s fast rebuilds mattered on laptop hardware.	Single-page delivery means a large first-load bundle, worsened by mathematical typesetting fonts.
KaTeX	Renders $\text{\LaTeX}$ mathematics synchronously and quickly, and outputs accessible markup rather than images [21].	Supports a subset of $\text{\LaTeX}$; unsupported macros in generated content must be caught before serving.
Tailwind CSS	Utility classes kept styling consistent across developers without a design system.	Verbose markup; class strings are hard to review in diffs.
Gemini API	After local generation proved infeasible on CPU (Section 2.4), a hosted model supplied the throughput; free-tier keys could be pooled to raise it further.	Rate limits shape the pipeline’s architecture; a hosted model is a dependency the project does not control, and model updates can silently change output quality.
Sentence-transformers	Multilingual embeddings are required for a Bangla corpus; English-only models retrieve poorly here [17].	A large download and CPU-bound indexing; embedding quality on Bangla mathematical text is materially worse than on English prose.
Git and GitHub	Version control and the integration point for five developers working on separable subsystems.	—

4.4 Implementation

Engine. The model was implemented from the equations rather than from a library, which was the right call for auditability: the policy document and the code state the same constants, and a reviewer can check one against the other. The graph enforces its own invariant — adding a

prerequisite edge provisionally inserts it, runs a three-colour depth-first search over the whole graph, and rolls the edge back if it would close a cycle. The graph cannot enter an invalid state through the public interface.

Ontology. Generated artefacts are never hand-edited. The source is compiled by script, and the editorial decisions that cannot be derived from data — topic aliases, display labels, course and section layout — live in a single configuration file. Rebuilding is a one-command operation, which is what made a late full rebuild affordable.

Content pipeline. Two failures shaped the implementation more than any design intent. The first: source books are machine-extracted, and one corrupt record in the Chemistry corpus — a string containing a $\LaTeX$ macro repeated tens of thousands of times and never terminated — desynchronised the parser and destroyed 334 neighbouring records. The parser was rewritten to resynchronise by decoding one object at a time, so a bad record now costs only itself. The second: language models emit $\LaTeX$ with single backslashes, and a strict JSON parser silently consumes several of those escapes as control characters rather than raising an error, corrupting text invisibly. This affected a measured 13.4% of one early batch. The fix is an escape-repair pass applied before every parse, in both the ingest and generation paths. Both failures are documented in the repository so that a future contributor does not rediscover them.

Frontend. The mathematics rendering component is the non-obvious piece: it splits mixed Bangla-and- $\LaTeX$ strings on delimiters, renders the mathematical segments through KaTeX, renders the rest as text, and normalises newline sequences that arrive literally from generated content.

Process. Work was carried out in a shared Git repository with subsystem-aligned branches merged into a main line — 93 commits across roughly five and a half months from five contributors. In place of a formal CI service, the project relies on a credential-free test harness that any developer can run in seconds before merging (Section 5.2), plus a chain of handover documents that record the state and rationale of each subsystem at the point it changed hands. Those documents are themselves an engineering artefact: they exist because the project spanned several months and five people, and knowledge that lived only in one contributor’s head was repeatedly found to be the project’s real bottleneck.

4.5 Deployment

Target environment. The system deploys as two processes and one managed service: a Python ASGI application serving the API, a static bundle for the client, and the hosted PostgreSQL instance. A container definition is provided for the backend. Configuration is by environment variables — the database URL and service key — supplied through a file that is excluded from

version control; the client takes the API base URL at build time and defaults to the local backend.

Release process. The ontology is compiled and its report inspected; generated questions are validated and loaded with a dry run first; the backend and client are then started. The deliberate ordering reflects the dependency chain in Figure 4.1: content cannot be loaded against an ontology that has not been compiled.

Operation and maintenance. The prototype is operated by the project team. Monitoring is limited to structured logging of question-search behaviour — each search prints which tier matched and how long the database query took, which is what made the topic-code defect in Section 5.3 visible. Backup is delegated to the managed database service.

Resources consumed. One small application process (well under a gigabyte of memory, since the loaded graph is a few thousand nodes), one static file bundle, and a database holding tens of megabytes. Offline, the pipeline consumes CPU for embedding the corpus and external API calls for generation — the only materially resource-hungry part of the system, and the subject of the environmental assessment in Section 6.4.

Chapter 5

Investigation and Evaluation

5.1 Experimental Investigation

This is a development-focused project, so the investigations reported here are those that validated — or refuted — specific design decisions. Four are presented. Each states the question it was asked to settle, the method, and the conclusion actually drawn, including the two cases where the investigation overturned the team’s initial position.

5.1.1 Investigation 1: is local model inference viable for content generation?

Question. Can the question bank be generated on the team’s own hardware using locally hosted open-weight models, avoiding any external API dependency and any cost?

Method. A generation pipeline was implemented against locally served quantised models (7B and 3B parameter classes) on CPU-only laptops, and measured on real specifications drawn from the ontology, with prompt sizes of roughly 4–6 k tokens and completions of roughly 3 k tokens per call.

Findings. Two independent failures. On throughput, the measured per-call latency extrapolated to multiple weeks of continuous computation for a single full pass over the 2,580-specification space then in scope — longer than the remaining project schedule, on machines also needed for development. On correctness, a subtler defect appeared: models emit $\text{\LaTeX}$ commands such as `\tan` and `\frac` with single backslashes, and a strict JSON parser does not reject these — for the subset of letters that happen to be legal JSON escapes it silently consumes the backslash and substitutes a control character. The corruption therefore produced no error at all, and was found only by reading output.

Conclusion. The local approach was abandoned on throughput grounds, and a hosted API adopted with key pooling to raise the rate ceiling. The escape defect was not specific to local models — it recurred with the hosted model at a measured 13.4% of one early production batch — so an escape-repair pass was made unconditional in both the ingest and generation paths. This is the clearest case in the project of an investigation changing the architecture: the design in

Section 4.2.9 exists because this one failed.

5.1.2 Investigation 2: is the ontology’s granularity sufficient to support the gating policy?

Question. The conjunctive gate in Equation (4.5) acts on a skill’s direct prerequisites. Does the first ontology supply enough structure for that gate to do anything?

Method. Structural analysis of the 430-skill graph: skills per topic, edges per skill, and the fraction of skills with no prerequisite edge at all.

Findings. The graph averaged fewer than four skills per topic and 1.1 edges per skill, with a substantial fraction of skills having no incoming edge. A gate that has no parents to test degenerates to the ungated case, and a topic holding three skills cannot localise a failure within itself — which is the system’s entire purpose.

Conclusion. The ontology, not the engine, was the binding constraint. This motivated the full-corpus rebuild (Section 4.2.6), which raised the graph to 1,688 skills and 1,743 edges with every one of the 183 syllabus topics populated and no isolated skills. Table 5.1 reports the rebuild’s structural effect.

Table 5.1: Structural effect of the ontology rebuild and its editorial pass

Property	Before editorial pass	After
Syllabus topics holding no skill	54	0
Unresolved prerequisite references	46	0
Isolated skills (no edge at all)	6	0
Topics with no internal learning order	5	0
Near-orphan fragments (2–6 skills)	91	32
Connected components	187	87
Largest connected component	203	482
Maximum DAG depth	7	9

5.1.3 Investigation 3: can prerequisite edges be proposed mechanically?

Question. After the rebuild, 91 fragments of 2–6 skills remained disconnected from the main graph. Can the missing edges be proposed automatically rather than authored by hand?

Method. Two scoring functions were implemented and their proposals read individually against the subject matter. The first scored candidate parents by token overlap with the orphaned skill. The second scored by *topic hub-ness* — preferring the candidate at the lowest Bloom tier with the most dependents.

Findings. Token overlap produced roughly 40% usable proposals. Its failures were systematic rather than random: it proposed lateral links between peers that merely share vocabulary — zero-order kinetics “requiring” first-order kinetics, the cosine rule “requiring” the sine rule — and in one case proposed a pair in both directions, which would have closed a cycle. Hub-based scoring raised usable proposals to roughly two thirds. Of 83 proposals finally reviewed, 48 were applied (several with the direction reversed from what was proposed) and 35 were rejected outright, their fragments left deliberately disconnected rather than wrongly connected.

Conclusion. Mechanical proposal is a useful accelerator and an unacceptable authority. A wrong prerequisite edge is worse than a missing one, because Equation (4.4) propagates along it and will credit a learner for a skill they have never demonstrated. The 32 fragments that survive are recorded as requiring a domain expert rather than another heuristic.

5.1.4 Investigation 4: do automated quality heuristics over-flag?

Question. Two validators were written to find defects in the generated ontology: one flagging skills that fuse two actions, and one flagging “Bloom inversions” — edges whose prerequisite sits at a *higher* Bloom tier than the skill requiring it. Are their outputs trustworthy?

Method. Every flag from both checks was read against the underlying skill or edge.

Findings. The multi-action check flagged 173 skills, of which over half were false positives, and was replaced by a stricter criterion yielding 47 candidates — of which 20 were genuine splits and 27 were rejected on reading. The Bloom-inversion check flagged 147 edges; hand review found 3 genuinely reversed. The check’s premise was wrong, not the data: Bloom tier measures cognitive demand, not teaching order, and a low-tier fact *about* an operation is routinely learned after the operation itself. For example, recalling that a determinant with two equal rows is zero (Remember) legitimately depends on being able to evaluate a determinant (Apply).

Conclusion. The inversion check was downgraded from error to advisory rather than deleted, with its reasoning recorded so the next contributor does not reinstate it. The transferable finding is that on this project mechanical heuristics over-flagged by roughly the same large factor twice, and that the correct response is to calibrate a heuristic against hand review before treating its output as a defect list.

5.2 Performance Evaluation

5.2.1 Setup and acceptance criteria

The engine is evaluated by an automated end-to-end harness that runs the real diagnostic, mastery-update, topic-practice and spillover code paths against an in-memory stand-in for the database, seeded with a synthetic learner and a question bank drawn from the compiled ontology. It requires no credentials and no network, which is what makes it runnable by any developer before a merge. Acceptance is binary per check: all checks must pass.

5.2.2 Results

The harness reports **18 checks passed, 0 failed**, reproduced on the current tree. Table 5.2 lists what each check establishes.

The ontology is separately validated by a structural checker, which reports **zero errors**: the graph is acyclic, every one of the 183 topics holds at least one skill, no skill is isolated, and no prerequisite reference is unresolved.

5.2.3 Evaluation against requirements

5.2.4 Evaluation against objectives

O1–O5, O7 and O8 are met as evidenced above. **O6 (question bank) is partially met**: the pipeline is built, validated and has produced and loaded real verified items, but full coverage of the specification space was not completed within the project schedule. **O3 is met in implementation but unvalidated in effect** — the propagation is provably consistent, but whether the resulting estimates correspond to real learner knowledge is exactly the question the next subsection says was not answered.

5.2.5 Limitations of this evaluation

The evaluation above is structural and functional. It establishes that the system computes what it was designed to compute and that its invariants hold. It does **not** establish that the system helps anyone learn, and the report makes no such claim. Specifically:

- **No human participants.** There was no user study, no classroom trial, and no usability testing with members of the target cohort. Every claim about learner benefit in this report

Table 5.2: End-to-end verification harness: checks and what each establishes

Area	Check	Establishes
Ontology	Graph builds (430 skills); catalogue topics all resolve (66); no legacy topic codes remain	The served ontology is internally consistent and the catalogue cannot reference a topic that does not exist
Question search	Tier-1 filter matches by topic <i>code</i> ; filtering by display label matches nothing	Regression guard for the defect in Section 5.3
Diagnostic	Session starts at length 30; question carries its subject; topic shown as label not code; all 30 questions answered	The adaptive loop runs to completion and returns learner-presentable metadata (FR3, FR12)
Mastery	Mastery unlocked after diagnostic; written to <code>user_skill</code> (43 skills above zero); section state reports subject	Evidence from 30 questions propagates to far more than 30 skills — the propagation mechanism doing its intended work (FR4, FR5, FR7)
Invariant	No ancestor sits below its descendant	Equation (4.4) holds globally after a full session — the correctness property in NFR1
Mastery view	View unlocked; graph returns 53 nodes and 52 edges; table rows carry subject	The visualisation contract holds (FR10)
Practice	Topic practice serves questions (8 answered)	Post-diagnostic flow works (FR8)
Spillover	Fires on repeated wrong answers, reporting rule B	The prerequisite-detection policy triggers on the intended evidence (FR9)

Table 5.3: Requirements traceability and status

Req.	Status	Evidence or deviation
FR1	Met, reduced	Account creation and sign-in work, but without a password (Section 5.3.3)
FR2	Met	Catalogue endpoint serves ontology-derived courses, sections and topics; verified by harness
FR3	Met	Diagnostic runs to a full 30-question session in the harness
FR4	Met	Bloom-conditioned update implemented per Table 4.2; exercised on every answered question
FR5	Met	Ancestor invariant verified explicitly; 30 answers moved 43 skills
FR6	Met	Conjunctive gate implemented per Equation (4.5)
FR7	Met	Mastery view locked before, unlocked after the diagnostic
FR8	Met	Topic practice served and answered questions in the harness
FR9	Met	Spillover fired and reported its triggering rule
FR10	Met	Mastery graph (53 nodes, 52 edges) and table both returned
FR11	Met	Retake endpoint resets section mastery and clears sessions
FR12	Met	Mathematics renders through KaTeX; literal newline handling fixed during implementation
FR13	Partially met	Pipeline is complete and validated, but the bank covers only part of the specification space (Section 5.2.5)
FR14	Met	Interface strings available in English and Bangla
NFR1	Met	Acyclicity enforced at insertion; ancestor invariant checked by harness; validator reports zero errors
NFR2	Not measured	No latency benchmark was run (Section 5.2.5)
NFR3	Met	Fallback ladder exercised; a missing item marks the skill tested and advances
NFR4	Met	Ontology compiled from source by script; generated files are not hand-edited
NFR5	Met	Harness runs with no credentials and no network
NFR6	Met	Entire system developed and run on CPU-only laptops
NFR7	Met	Bangla content with L ^A T _E X mathematics end to end
NFR8	Met	Credentials held in a git-ignored environment file, server-side only

is an argument from design, not a measurement.

- **No fitted parameters.** $P(L_0)$, $P(T)$ and the Bloom-conditioned $P(G)$, $P(S)$ values are pedagogical priors, not estimates. The literature is explicit that fitted and individualised parameters outperform fixed ones [4, 5]; this system cannot do that until it has response data, which it cannot have before deployment.
- **No predictive validation.** The standard quantitative check on a knowledge-tracing model — does it predict the next response better than a baseline, measured by AUC on held-out data — was not performed, because there is no held-out learner data.
- **No performance benchmark.** NFR2 was not measured; the system is subjectively responsive in use, which is not evidence.
- **Ontology unreviewed by domain experts.** The rebuilt graph is structurally valid but its subject-matter correctness has been verified only by the team. 32 disconnected fragments and a body of hand-authored prerequisite edges are explicitly awaiting review by a chemist and a physicist.
- **Question bank coverage is partial,** so a full-coverage session is not yet reachable for every topic.

5.3 Deviations and Design Revisions

5.3.1 Topic filter matched no rows (found, analysed, fixed)

Deviation. The diagnostic’s first two search tiers returned nothing, ever. Every question was silently arriving from the untargeted fallback, so topic-targeted and graph-nearby selection — two thirds of the selection design — were dead code in practice.

Cause. The `questions.topic` column is a foreign key to `ontology_topics` and therefore holds a topic *code* (MAT_MATRIX), while the search filtered it by the catalogue *display label* (“Matrices and Determinants”). The two can never be equal, so the filter matched no row. Nothing failed loudly: a query returning zero rows is a valid query, and the fallback dutifully supplied a question, so sessions looked healthy.

Revision. The filter was corrected to compare codes, and — more importantly — two checks were added to the harness: one asserting that a code-based filter matches, and one asserting that a label-based filter matches *nothing*. The second check is the one that matters; it fails if anyone reintroduces the confusion.

5.3.2 Session state is not persistent

Deviation. Diagnostic and practice sessions live in process memory, so a backend restart discards every in-flight session (NFR-adjacent, and a production defect).

Analysis. The session objects hold graph traversal state that was not designed for serialisation. Persisting them properly means either a session store or reconstructing session state from the durable mastery table on demand.

Decision. Not revised, deliberately. The deployment target for this project is a single demonstration process where the failure mode is a lost session rather than lost mastery — mastery itself *is* persisted. The correct fix is recorded as production hardening rather than disguised.

5.3.3 Authentication is username-only

Deviation. A learner signs in with a username and no password; the client holds the resulting identity in browser storage. This does not meet the access-control expectations of Section 3.2.

Analysis. The system stores no data that is sensitive in itself beyond performance history, and the prototype's users are the team. Authentication was therefore deprioritised against the ontology work, which was on the critical path for every other subsystem.

Decision. Not revised within the project. It is the first item of production hardening (Section 7.1.3), and is stated here rather than presented as a design choice, because it is not one — it is an accepted debt.

5.3.4 The rebuilt ontology is not yet the served ontology

Deviation. The application serves the 430-skill legacy ontology; the validated 1,688-skill rebuild is complete but not adopted.

Analysis. Adoption is not merely a file swap. The question bank is keyed to legacy skill identifiers, so adopting the new graph without regenerating content would leave most new skills with no questions — a system that knows about 1,688 skills but can ask about few of them is worse for a learner than one that knows about 430 and can ask about most.

Decision. Deliberately sequenced rather than rushed: adopt after a budgeted generation run against the new specification space (Section 6.1.2). The dependency is recorded so the sequence is not attempted in the wrong order.

Chapter 6

Project Management, Teamwork, and Impact

6.1 Project Management and Finance

6.1.1 Planning and Risk Management

The project ran across one academic session with five contributors, all carrying a full course load. Work was partitioned along the subsystem boundaries of Figure 4.1, which was chosen partly for architectural reasons and partly because those boundaries are the ones that let several people work without blocking each other: the engine, the content pipeline and the client each depend on the ontology’s *interface* but not on each other’s internals.

Figure 6.1 shows the schedule. The version-control history records 93 commits between April and September 2026, and the phases below are drawn from it rather than from an idealised plan.

Where actual progress differed from plan. Three deviations, each instructive.

1. **A late, unplanned ontology rebuild.** The plan assumed the knowledge structure was a one-time input. It turned out to be the project’s critical path: the first ontology was too sparse for the gating policy to act on, and rebuilding it consumed most of September. The schedule absorbed this only because the compiler-based build made a rebuild a one-command operation.
2. **A discarded content-generation approach.** Weeks of work on local model inference were written off when measurement showed it could not finish in the remaining schedule (Investigation 1). Migrating to a hosted API recovered the capability but consumed schedule that had been allocated to coverage.
3. **Scope reduction on authentication.** Deferred deliberately once it became clear the ontology work would not otherwise finish.

Table 6.1: Risk register: principal risks, their realisation, and mitigation

Risk	Realised?	Mitigation and outcome
Content generation infeasible at required scale	Yes	Measured early rather than assumed; switched to a hosted API with pooled keys. Cost the schedule, not the deliverable.
Knowledge structure inadequate for the inference built on it	Yes	Detected by structural analysis rather than by a user complaint; rebuild was affordable because the ontology build was automated from the start.
Machine-generated questions are wrong or off-topic	Partly	Two-model generate-then-verify with blind re-solve; only items passing both checks are written.
Source corpus corruption propagating silently	Yes	One malformed record destroyed 334 neighbours before the parser was made resynchronising; escape repair made unconditional.
External API rate limits or cost overrun	Yes	Key pooling, adaptive backoff, batch verification, resumable progress files. One parallelised run was halted mid-flight when its consumption rate proved unsustainable, and the work was restructured to a cheaper sequential method.
Knowledge loss across a five-person, multi-month project	Partly	A chain of handover documents recording state and rationale per subsystem; these were written precisely because context kept being lost between sessions and contributors.
Team member unavailability during examinations	Partly	Subsystem boundaries chosen so that no single absence blocked others.
Secret leakage into a public repository	No	Credentials confined to a git-ignored environment file from the first commit that needed one; never committed.




Figure 6.1: Project schedule. The ontology rebuild in September was not in the original plan; it was forced by Investigation 2 (Section 5.1).

6.1.2 Budgeting and Resource Identification

The project was built entirely on free tiers and hardware the team already owned; direct monetary expenditure was zero. That is not the same as costless, so Table 6.2 itemises both the actual outlay and the notional value of what was consumed, and Table 6.3 models what a real deployment would cost.

The free tiers are a genuine constraint rather than a lucky break: the pipeline’s key-pooling, adaptive backoff and resumable-progress design exist specifically because per-key rate limits, not money, were the scarce resource.

6.1.3 Economic Analysis and Financial Prospecting

The cost structure has one useful property: it is almost entirely fixed. Content generation and ontology construction are one-time costs; serving an additional learner costs a database row and a few constant-time updates. The marginal cost per learner is therefore close to zero, which is the necessary condition for any viable offering in this market, because Section 2.3 established that individual willingness to pay for software is low even though willingness to pay for preparation is high.

Three revenue paths follow from that, in descending order of realism:

Table 6.2: Development resources: actual and notional cost

Resource	Basis	Actual cost
Engineering effort	5 students, one academic session, part-time alongside coursework	BDT 0 (in kind)
Development hardware	Personal laptops, CPU only, no GPU purchased	BDT 0 (owned)
Database and hosting	Supabase free tier	BDT 0
Model API for generation	Gemini free-tier keys, pooled across accounts to raise the rate ceiling	BDT 0
Embedding model	Open-weight multilingual model, run locally	BDT 0
Source content	Textbooks and question banks already owned by team members	BDT 0
Version control and CI	GitHub free tier	BDT 0
Total direct expenditure		BDT 0

Table 6.3: Indicative operating cost for a production deployment (order-of-magnitude estimate, not a quotation)

Item	Driver	Scale
One-time question generation	1,688 skills $\times$ 6 Bloom levels = 10,128 specifications, each a generation call plus a share of a batched verification call	Dominant one-time cost
Managed database	Tens of MB of content; one row per learner per skill	Lowest paid tier adequate at pilot scale
Application hosting	One small always-on process; inference is constant-time per response	Smallest instance class
Content refresh	Regeneration only when the syllabus or ontology changes	Infrequent
Domain expert review	Chemist and physicist review of the ontology (Section 5.2.5)	The largest genuinely unavoidable human cost

1. **Institutional licensing.** Coaching centres and schools pay per cohort for diagnostic reporting that they cannot produce manually. The buyer has budget, the value proposition is diagnosis at a granularity a human cannot reach across a large class, and the number of buyers is small enough to reach directly.
2. **Freemium to learners.** Diagnostic free, unlimited practice paid. This fits observed behaviour but yields little per user and depends on scale the project does not have.
3. **Public or NGO deployment.** Offered at no cost as an equity intervention (Section 6.5.1), funded by a body whose objective is access rather than revenue. This is the path most consistent with the system's social value, and the least dependent on a commercial market.

The honest financial reading is that this is not, in its current form, a business: it is a prototype whose economics work only after the one-time content cost is paid and whose evidence of efficacy — the thing an institution would actually buy — does not yet exist (Section 5.2.5). Establishing that evidence is a prerequisite for any of the three paths above, not a follow-on.

6.1.4 Sustainability in Business and Commercialisation

For the system to outlive the course, three things must be true. First, the knowledge structure must be maintainable by someone who is not the original team: this is addressed architecturally, since the ontology is compiled from a source file by a documented one-command build, with editorial decisions isolated in a single configuration file. Second, the content must be regenerable rather than hand-curated, which the pipeline provides. Third, the system must be transferable — which is why the repository carries an operational guide and a chain of handover documents written for a reader with no prior context, and why the test harness runs without credentials.

The realistic continuation path is academic rather than commercial: the codebase, the validated ontology and the documented open questions form a defined starting point for a subsequent capstone or thesis, with the domain expert review and a controlled efficacy study as the natural next pieces of work. A commercial path would require a sponsor to fund the generation run and the expert review before any user-facing launch.

6.2 Individual and Teamwork

6.2.1 Participation and Contribution

Ownership was primary rather than exclusive — the ontology rebuild in particular drew in contributors from outside its owning area, because it sat on every other subsystem's critical path.

Table 6.4: Contribution by subsystem

Member	Primary area	Contributions
Sayaad Muzahid Masfi (2105066)	Adaptive engine	BKT implementation, the three-phase mastery policy, the skill-graph data structure with its cycle check, and the policy specification document that the code is audited against.
Aurchi Chowdhury (2105083)	Backend API & sessions	FastAPI service, session lifecycle for diagnostic and practice, the tiered question search, and the spillover rules.
Omar Abdur Razzaque (2105094)	Content pipeline	Corpus parsing and repair, retrieval index, the generate-then-verify generation pipeline, and the database loader with its validation pass.
Sakif Naieb Raiyan (2105065)	Frontend	React client, the question-answer flow, mastery visualisation, bilingual interface, and mathematical rendering.
Rubaiyat -E-Zaman (I2105104)	Ontology	Skill extraction from the corpus, the ontology compiler, the full-corpus rebuild and its editorial pass, and the structural validators.

6.2.2 Collaboration and Conflict Resolution

Collaboration was organised around the repository. Work proceeded on subsystem-aligned branches merged into a main line, so that the integration point was explicit and reviewable rather than continuous. The credential-free test harness served as the shared definition of “not broken”: any contributor could run it in seconds before merging, which removed most integration disagreements by making them checkable rather than arguable.

The substantive technical disagreements in the project were about the ontology — specifically about granularity, since finer skills improve diagnosis but multiply content cost, and different contributors weighted those differently. These were resolved by writing the criteria down and testing them against data: the decision to rebuild was settled by the structural analysis in Investigation 2 rather than by argument, and the granularity target was fixed by reference to what the gating policy needs to function. The general pattern the team converged on was to convert a disagreement about judgement into a measurement wherever one could be defined.

6.2.3 Leadership and Direction

Direction was distributed by subsystem rather than centralised: each area’s owner set its technical direction and was accountable for it, with cross-cutting decisions — the model family, the granularity target, the decision to rebuild — taken jointly. The mechanism that kept direction coherent across a five-and-a-half-month project was documentary: each significant change was

accompanied by a handover note recording not just what changed but why, so that a contributor returning to an area after weeks away, or a new contributor entering it, inherited the reasoning rather than only the code.

The clearest instance of planning after a setback is Investigation 1. When local generation was measured and found unable to finish within the schedule, the response was not to persist but to re-plan: the approach was abandoned, the alternative adopted, and the reasoning recorded so that the abandoned path would not be re-attempted by someone who had not seen the measurements.

6.2.4 Multidisciplinary Engagement

The project is multidisciplinary by construction: it required computer science (inference, graph algorithms, systems), educational psychology (knowledge tracing, Bloom's taxonomy, the zone of proximal development), and subject-matter expertise in three natural sciences at HSC level. The team's engagement outside its own discipline took three forms: reading and applying the education research literature rather than treating pedagogy as common sense; working directly with HSC Mathematics, Physics and Chemistry source material to build the ontology; and explicitly identifying the boundary of its own competence — the decision to mark 32 graph fragments and a body of authored chemistry edges as requiring a chemist's and a physicist's review, rather than resolving them by a computer scientist's intuition, is the clearest example of recognising where another discipline's authority was needed.

6.3 Impact on Society

Positive: access. The dominant positive impact is distributional. Expert diagnostic attention is currently rationed by money and geography — the best coaching is concentrated in cities and priced accordingly. A system whose marginal cost per learner is near zero can offer a form of that attention to a student in a district town with a phone and a connection. This is the project's principal social justification.

Positive: efficiency of effort. Redirecting a student from repeated failure on a downstream topic to the upstream skill actually missing reduces wasted study time in a year where time is the binding constraint.

Negative: misdiagnosis. The corresponding harm is specific and worth stating plainly. A wrong prerequisite edge or a miscalibrated parameter can send a learner to study something they already know, or release them from something they do not. In a high-stakes year that is a real cost borne by someone with no way to detect the error. The design response is the conservative mastery threshold and the refusal to guess prerequisite edges mechanically (Investigation 3); the honest

response is that this risk is why efficacy evidence is a precondition for deployment, not a nicety.

Negative: intensification. A system that always has a next question can extend an already unhealthy preparation culture. The design does not currently counter this — there are no session limits or break prompts — and adding them is listed in future work.

Health and safety. No physical risk. The relevant health dimension is psychological: adaptive systems that target a constant challenge level deliberately keep the learner near their failure boundary, which is motivationally corrosive if presented as bare failure. Mitigation in the current design is that mastery is shown as a graph of what has been secured, not only as what remains.

Cultural. Bangla-first content with correct mathematical typesetting is a deliberate cultural position: it refuses the common assumption that serious technical study must happen in English, and it serves students whose mathematical vocabulary is Bangla.

Legal. Two obligations govern deployment: personal data of a predominantly under-20 user base, and copyright in source material. Both are addressed in Section 3.2; the current prototype's lack of authentication means it is not yet lawful to deploy publicly with real learner data, which is stated as a blocking condition rather than a limitation.

6.4 Environmental Impact and Life Cycle Sustainability

Development. Modest: five laptops over one session, and CPU-bound embedding of a corpus of a few thousand records. The abandoned local-inference approach would have been materially worse — weeks of continuous CPU computation — so the decision taken on schedule grounds in Investigation 1 also reduced energy consumption.

Content generation. This is the project's environmental hotspot. Large language model inference carries a real energy and water footprint in the provider's data centre, and the specification space at rebuild scale is 10,128 generation calls plus batched verification. Three design choices reduce it, all originally made for cost or throughput reasons: verification uses a smaller model than generation; verification calls are batched across several specifications; and progress is checkpointed so an interrupted run resumes rather than restarting. A fourth mitigation is intrinsic — generation is a one-time cost, not a per-learner one, and it amortises over every learner thereafter.

Operation. Very low. One small always-on process and a managed database; inference per response is arithmetic, not model inference. The system does no work while a learner is reading a question.

Retirement. Two obligations. Learner data must be deletable, which the schema supports through cascading deletion from the learner record, and which a production deployment should expose as

an explicit account-deletion path. The generated question bank and ontology are the project's durable artefacts and are worth preserving beyond the system itself — as a syllabus-aligned skill graph they have value independent of this application, and releasing them rather than discarding them is the sustainable end-of-life choice.

6.5 Ethics

6.5.1 Equity and Inclusivity

Language. The system is Bangla-first for content, with a bilingual interface. This is the single largest inclusivity decision in the project: it admits students whose technical vocabulary is Bangla, who are systematically disadvantaged by English-only tools.

Cost. Zero marginal cost per learner makes free provision technically possible, which is the precondition for reaching students for whom coaching fees are prohibitive.

Access assumptions, and who is still excluded. The design assumes a browser and a connection. That excludes students without either — which is precisely the group the access argument above claims to serve. This tension is not resolved in the current system. Offline capability and low-bandwidth operation are the honest mitigations, and neither is implemented.

Disability. Mathematics is rendered as accessible markup rather than as images, and controls are keyboard-operable, but no formal accessibility audit against WCAG AA [21] was performed and no screen-reader testing was done. Stating this as an outstanding gap is more useful than claiming compliance that has not been verified.

Fairness across learners. One model-level equity issue deserves naming: the parameters in Table 4.2 are uniform across all learners. A learner whose true guess or slip behaviour differs from those priors — for instance, a careful learner who slips rarely — is modelled slightly wrongly, and uniformly. Individualised parameters [5] are the known remedy and require response data the system does not yet have.

6.5.2 Accountability and Personal Responsibility

Responsibility for the system's behaviour rests with the team, and within the team it is traceable: subsystem ownership is recorded (Table 6.4) and every change is attributable through version control. The project's accountability practices were concrete rather than declarative — the defect in Section 5.3 was reported, analysed and guarded against rather than quietly fixed, and the two over-flagging validators in Investigation 4 were documented as team errors of judgement rather than deleted.

The accountability commitment that matters most for this system is about claims. The team is responsible for not asserting that ATLAS improves examination outcomes, because no evidence for that exists. That commitment is why Section 5.2.5 exists in this report in the form it does.

6.5.3 Proper Use of Intellectual Property

Prior art and literature. The model family, the taxonomy and the retrieval approach are not original to this project and are cited to their sources [2, 3, 11, 13, 14, 16]. The original contribution claimed is the composition — Bloom-conditioned BKT with ancestor pull-up and conjunctive gating over a syllabus-derived graph — not its components.

Third-party software. All third-party components are open-source under permissive licences and are used as dependencies rather than copied into the codebase: FastAPI, Uvicorn, Pydantic (backend); React, Vite, Tailwind CSS, KaTeX, Lucide (frontend); sentence-transformers and NumPy (pipeline); PostgreSQL via Supabase. Their licences are carried in the dependency manifests.

Source content. The corpus is drawn from copyrighted HSC textbooks and published question banks. It is used as a private retrieval and grounding corpus; it is not redistributed, is excluded from distributable artefacts, and generated items are new text rather than reproductions. Any public deployment would require this position to be reviewed with the rights holders.

Generative AI disclosure. Generative AI was used substantively in this project and is disclosed accordingly:

- **As a system component.** The question bank is machine-generated via a hosted large language model, with a second model verifying each item. This is not an incidental use — it is a designed part of the architecture (Section 4.2.9), and the verification stage exists precisely because unverified model output is not trustworthy as assessment content.
- **As a development tool.** AI coding assistants were used during development for code generation, refactoring, documentation, the ontology extraction and consolidation passes, and in the preparation of this report. All such output was reviewed by team members before being committed, and the team accepts full responsibility for it (Section 6.5.2). Investigation 4 is a direct record of that review process catching machine-produced errors — both over-flagging validators were AI-assisted artefacts corrected by hand review.

6.5.4 Professionalism and Ethical Codes

The ACM Code of Ethics [23] and the IEEE Code of Ethics [22] both place avoiding harm and honest representation of work above the practitioner’s own interest. Three project decisions

follow directly from those obligations and were in tension with the team's interest in presenting a polished result:

1. **Refusing an unevidenced efficacy claim** (ACM 1.3, “be honest and trustworthy”). The straightforward way to strengthen this report would be to assert that the system improves learning. Section 5.2.5 instead states that no such evidence was gathered.
2. **Declining to connect prerequisite edges mechanically** (ACM 1.2, “avoid harm”). Accepting all 83 machine proposals would have produced a better-looking connectivity statistic. Thirty-five were rejected because a wrong edge silently corrupts a learner's diagnosis, which is a harm the statistic would have concealed.
3. **Documenting defects rather than quietly closing them** (IEEE 7.8.7, accepting and offering honest criticism of technical work). The topic-code defect, the corpus corruption and the over-flagging validators are all recorded in the repository and in this report with their causes.

The public-interest dimension specific to this system is that its users are young people making consequential decisions under pressure, largely without the expertise to audit the advice they are given. That asymmetry is the reason the team treats “do not claim more than the evidence supports” as the project's governing professional obligation rather than as a reporting formality.

Chapter 7

Conclusion

7.1 Conclusion and Future Work

7.1.1 Achievement against objectives

Table 7.1: Achievement against the objectives of Section 1.4

Obj.	Status	Basis
O1	Achieved	1,688 skills, 1,743 edges, all 183 syllabus topics populated; validated acyclic, depth 9, zero isolated skills, zero unresolved references
O2	Achieved	Bloom-conditioned BKT implemented from the equations; exercised on every response by the verification harness
O3	Achieved in implementation	Ancestor pull-up and conjunctive gating implemented; the ancestor-ordering invariant is checked explicitly and holds. Whether the resulting estimates match real learner knowledge is untested
O4	Achieved	A 30-question diagnostic completes and moves mastery on 43 skills — more skills than questions asked, which is the propagation doing its work
O5	Achieved	Topic practice runs to the mastery threshold; spillover fires on the intended evidence
O6	Partially achieved	Pipeline built, validated, and producing verified items; coverage of the full specification space incomplete
O7	Achieved	Seven-route web client with diagnostic, practice and mastery views, bilingual, with Bangla and L ^A T _E X rendering
O8	Achieved	18-check end-to-end harness, credential-free, passing completely

The project set out to build a tutor that could tell a student *which* foundation is missing rather than merely that they are wrong, and the mechanism for that — an interpretable per-skill model, propagating along an explicit prerequisite graph, gated conjunctively — is built and demonstrably self-consistent. The single most valuable thing the project learned is not in the engine at all: it is that the knowledge structure, not the inference algorithm, is the binding constraint on a system like this. A sophisticated update rule over a graph with 1.1 edges per skill has almost nothing to reason with. Most of the engineering effort in the second half of the project went into the ontology for exactly that reason, and that is the finding most likely to transfer to anyone

attempting similar work.

7.1.2 Limitations

Stated plainly, and in the order that matters:

1. **No efficacy evidence.** The system has never been used by a real learner. Every claim about learning benefit is an argument from design.
2. **No fitted parameters.** All model parameters are pedagogical priors; the literature is clear that fitted and individualised parameters do better [4, 5].
3. **The validated ontology is not the served ontology.** The application still runs on the 430-skill legacy graph, because adopting the rebuild without regenerating content would leave most skills unquestionable.
4. **Question-bank coverage is partial.**
5. **The ontology has not been reviewed by domain experts** — 32 disconnected fragments and a body of hand-authored chemistry edges await a chemist and a physicist.
6. **Production-readiness gaps:** no password authentication, in-memory session state, wide-open CORS, no accessibility audit, no performance benchmark.

7.1.3 Future work

Immediate, and blocking deployment. Password authentication and server-side sessions; persistent session state; CORS restriction. These are small, well-understood, and are the difference between a demonstrator and something that may lawfully hold real learner data.

Completing the knowledge base. A budgeted generation run against the $1,688 \times 6$ specification space, followed by adoption of the rebuilt ontology in the correct order (generate, then switch); domain-expert review of the 32 fragments and the authored edges.

Establishing efficacy — the most important item. A controlled study with real candidates: does the system's mastery estimate predict performance on held-out questions better than a baseline, measured by AUC as the knowledge tracing literature does; and does directing study by the system's diagnosis produce better outcomes than self-directed study? Until this is done, the central claim of the project is unverified.

Model improvements once data exists. Fitted and individualised BKT parameters [5]; contextual estimation of slip and guess [4]; and, with sufficient data, a comparison against a deep knowledge tracing baseline [10]. The engine was deliberately built with a stateless, replaceable model interface so that this substitution requires no change to its call sites.

Pedagogical extensions. Forgetting/decay in the mastery estimate, which the current model omits entirely and which matters over a months-long preparation period [18]; spaced review scheduling; and session pacing that resists the intensification risk noted in Section 6.3.

7.2 Continuous Engagement and Self-directed Learning

Little of what this project required was covered by prior coursework, and the learning it forced was mostly self-directed.

Domain literature learned independently. The knowledge-tracing literature was new to the team: the original BKT formulation [2], the variants that refine its noise parameters and individualise them [4–6], the competing model families [8–10], the prerequisite-structured student models closest to this design [7], and the older knowledge-space theory that supplies the graph formulation [11, 12]. The educational side — the revised Bloom taxonomy [3] and the zone of proximal development [14] — was read as source material for design decisions rather than as background, since both appear directly in the implemented policy.

Emerging technical practice. Retrieval-augmented generation [16] and multilingual sentence embeddings [17] were adopted from current practice; the survey literature on automatic question generation [15] was what convinced the team that generation without an independent verification stage would not produce usable assessment items, which directly shaped Section 4.2.9.

Tools and technologies learned during the project. FastAPI and Pydantic; React with Vite and Tailwind; Supabase and its constraint model; KaTeX and the specific difficulties of mixing Bangla text with $\text{\LaTeX}$ mathematics; local and hosted large-language-model inference including quantised local serving, and the practical engineering around hosted APIs — rate limiting, key pooling, adaptive backoff and resumable batch jobs.

Learning from failure, which was the majority of it. Three episodes taught more than any tutorial. Measuring local inference before committing to it would have saved weeks — the lesson being to measure the throughput of a critical dependency before building on it. The silent JSON escape corruption taught that a parser accepting input is not evidence the input was understood, and that data corruption which raises no error is the expensive kind. The two over-flagging validators taught that a heuristic must be calibrated against hand review before its output is treated as a defect list.

Practices adopted. The habit that most changed how the team works is documentary: writing down the reasoning behind a decision at the moment it is taken, not afterwards. On a five-person project spanning months, the recurring bottleneck was not writing code but recovering context — why a constant has the value it has, why an approach was abandoned, what a validator’s output actually means. The chain of handover documents in the repository exists because of that, and it

is the practice the team expects to carry into work beyond this course.

Bibliography

- [1] K. VanLehn, “The relative effectiveness of human tutoring, intelligent tutoring systems, and other tutoring systems,” *Educational Psychologist*, vol. 46, no. 4, pp. 197–221, 2011.
- [2] A. T. Corbett and J. R. Anderson, “Knowledge tracing: Modeling the acquisition of procedural knowledge,” *User Modeling and User-Adapted Interaction*, vol. 4, no. 4, pp. 253–278, 1995.
- [3] L. W. Anderson and D. R. Krathwohl, *A Taxonomy for Learning, Teaching, and Assessing: A Revision of Bloom’s Taxonomy of Educational Objectives*. Longman, 2001.
- [4] R. S. J. d. Baker, A. T. Corbett, and V. Aleven, “More accurate student modeling through contextual estimation of slip and guess probabilities in bayesian knowledge tracing,” in *Intelligent Tutoring Systems (ITS)*, 2008, pp. 406–415.
- [5] M. V. Yudelson, K. R. Koedinger, and G. J. Gordon, “Individualized bayesian knowledge tracing models,” in *Artificial Intelligence in Education (AIED)*, 2013, pp. 171–180.
- [6] Z. A. Pardos and N. T. Heffernan, “Kt-idem: Introducing item difficulty to the knowledge tracing model,” in *User Modeling, Adaptation and Personalization (UMAP)*, 2011, pp. 243–254.
- [7] T. Käser, S. Klingler, A. G. Schwing, and M. Gross, “Dynamic bayesian networks for student modeling,” *IEEE Transactions on Learning Technologies*, vol. 10, no. 4, pp. 450–462, 2017.
- [8] H. Cen, K. Koedinger, and B. Junker, “Learning factors analysis – a general method for cognitive model evaluation and improvement,” in *Intelligent Tutoring Systems (ITS)*, 2006, pp. 164–175.
- [9] P. I. Pavlik, H. Cen, and K. R. Koedinger, “Performance factors analysis – a new alternative to knowledge tracing,” in *Artificial Intelligence in Education (AIED)*, 2009, pp. 531–538.
- [10] C. Piech, J. Bassen, J. Huang, S. Ganguli, M. Sahami, L. J. Guibas, and J. Sohl-Dickstein, “Deep knowledge tracing,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2015, pp. 505–513.
- [11] J.-P. Doignon and J.-C. Falmagne, “Spaces for the assessment of knowledge,” *International Journal of Man-Machine Studies*, vol. 23, no. 2, pp. 175–196, 1985.

- [12] J.-C. Falmagne and J.-P. Doignon, *Learning Spaces: Interdisciplinary Applied Mathematics*. Springer, 2011.
- [13] B. S. Bloom, *Taxonomy of Educational Objectives: The Classification of Educational Goals, Handbook I: Cognitive Domain*. David McKay Company, 1956.
- [14] L. S. Vygotsky, *Mind in Society: The Development of Higher Psychological Processes*. Harvard University Press, 1978.
- [15] G. Kurdi, J. Leo, B. Parsia, U. Sattler, and S. Al-Emari, “A systematic review of automatic question generation for educational purposes,” *International Journal of Artificial Intelligence in Education*, vol. 30, no. 1, pp. 121–204, 2020.
- [16] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W.-t. Yih, T. Rocktäschel, S. Riedel, and D. Kiela, “Retrieval-augmented generation for knowledge-intensive nlp tasks,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [17] N. Reimers and I. Gurevych, “Sentence-bert: Sentence embeddings using siamese bert-networks,” *Proceedings of EMNLP-IJCNLP*, pp. 3982–3992, 2019.
- [18] H. Ebbinghaus, *Memory: A Contribution to Experimental Psychology*. Teachers College, Columbia University, 1885.
- [19] Government of the People’s Republic of Bangladesh, “Digital security and data protection policy framework,” 2023, tODO: verify the exact instrument and year applicable at submission time.
- [20] OWASP Foundation, “Owasp top 10: The ten most critical web application security risks,” 2021.
- [21] World Wide Web Consortium, “Web content accessibility guidelines (wcag) 2.1,” 2018, w3C Recommendation.
- [22] Institute of Electrical and Electronics Engineers, “Ieee code of ethics,” 2020.
- [23] Association for Computing Machinery, “Acm code of ethics and professional conduct,” 2018.